

in **binary classification**, your model is deciding between **two options** — usually called **Positive** and **Negative** (like “spam” vs. “not spam”).

Precision

Precision answers:

“When my model says *Positive*, how often is it actually right?”

It’s about **how accurate your positive predictions are**.

Example:

If your spam filter marks 10 emails as spam and 8 really are spam,

👉 Precision = $8 / 10 = 0.8$ (80%)

So, high precision means **few false alarms**.

Recall

Recall answers:

“Out of all the real *Positives*, how many did my model find?”

It’s about **how many true positives you caught**.

Example:

If there are 20 spam emails total and your filter caught 8,

👉 Recall = $8 / 20 = 0.4$ (40%)



So, high recall means **you found most of the actual positives**.

The Trade-Off

- If you **raise precision**, you become more cautious (fewer false positives) but might **miss some true ones** → lower recall.
 - If you **raise recall**, you catch more positives but might **get more false alarms** → lower precision.
-

Quick Analogy

Imagine fishing:

-  **Recall** = how many total fish you catch out of all the fish in the lake.
-  **Precision** = how many things in your net are actually fish (not boots or cans).